

1. Project Summary

In this notebook, we aim to analyze data from an aviation data set from National Transport and Safety Board.

This data contain accident information from 1962 - 2023.

2. Business Understanding

We will use the above dataset to determine for our company which is the safest airplane to purchase.

The data set contains several columns but we will focus on those that relate to airplane safety.

3. Data Understanding -Data Source

This data source was sourced from NTSB Aviation Accident Database and involves data concerning accidents of different aircraft makes between 1962- 2023. For this project despite the many columns of data,we will use the total fatal injuries column and engine types column to suggest the best aircraft make.

4. Business Objectives

This project aims at helping my company make a decision on which aircraft have the lowest risk to purchase and operate by giving clear, actionable recommendations

5. Success Criteria

We will only suggest aircrafts with the lowest fatal injuries and most effective engine type.

6. Data Preparation

6.1 Data Loading

```
In [1]: #Loading Data and importing necessary libraries  
# import pandas and matplotlib  
  
import pandas as pd  
import matplotlib.pyplot as plt  
%matplotlib inline
```

In [6]: # Let us read our data set

```
aviation_data = pd.read_csv('Aviation_Data.csv')
```

c:\Users\Moringa School\anaconda3\envs\learn-env\lib\site-packages\IPython\core\interactiveshell.py:3145: DtypeWarning: Columns (6,7,28) have mixed types.Specify dtype option on import or set low_memory=False.

```
has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```


In [7]: # let us confirm that our data has been read by looking at the first 5 rows

```
aviation_data.head()
```

Out[7]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States	NaN	NaN
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States	NaN	NaN
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States	36.9222	NaN
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States	NaN	NaN
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States	NaN	NaN

5 rows × 31 columns



6.2 Data Exploration

In [8]: # Looking at our columns

```
aviation_data.columns
```

Out[8]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date', 'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code', 'Airport.Name', 'Injury.Severity', 'Aircraft.damage', 'Aircraft.Category', 'Registration.Number', 'Make', 'Model', 'Amateur.Built', 'Number.ofEngines', 'Engine.Type', 'FAR.Description', 'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status', 'Publication.Date'],
dtype='object')

In [9]: # Let us explore our data set generally

```
aviation_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 90348 entries, 0 to 90347
Data columns (total 31 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Event.Id        88889 non-null   object  
 1   Investigation.Type  90348 non-null   object  
 2   Accident.Number  88889 non-null   object  
 3   Event.Date       88889 non-null   object
```

```

4   Location          88837 non-null  object
5   Country           88663 non-null  object
6   Latitude          34382 non-null  object
7   Longitude         34373 non-null  object
8   Airport.Code      50249 non-null  object
9   Airport.Name      52790 non-null  object
10  Injury.Severity  87889 non-null  object
11  Aircraft.damage   85695 non-null  object
12  Aircraft.Category 32287 non-null  object
13  Registration.Number 87572 non-null  object
14  Make              88826 non-null  object
15  Model              88797 non-null  object
16  Amateur.Built     88787 non-null  object
17  Number.of.Engines  82805 non-null  float64
18  Engine.Type        81812 non-null  object
19  FAR.Description    32023 non-null  object
20  Schedule           12582 non-null  object
21  Purpose.of.flight  82697 non-null  object
22  Air.carrier        16648 non-null  object
23  Total.Fatal.Injuries 77488 non-null  float64
24  Total.Serious.Injuries 76379 non-null  float64
25  Total.Minor.Injuries 76956 non-null  float64
26  Total.Uninjured    82977 non-null  float64
27  Weather.Condition  84397 non-null  object
28  Broad.phase.of.flight 61724 non-null  object
29  Report.Status      82508 non-null  object
30  Publication.Date   73659 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.4+ MB

```

In [10]: *## Selecting the columns we will need for our analysis*

```
selected_columns = ['Total.Fatal.Injuries', 'Make', 'Engine.Type']
```

In [11]: *# Lets us restrict our data set to the pre-defined selected columns*

```
aviation_data = aviation_data[selected_columns]
```

In [12]: *# Let us confirm that our data set now contains information of only the 2 columns
We will read the first 5 rows.*

```
aviation_data.head()
```

Out[12]:

	Total.Fatal.Injuries	Make	Engine.Type
0	2.0	Stinson	Reciprocating
1	4.0	Piper	Reciprocating
2	3.0	Cessna	Reciprocating
3	2.0	Rockwell	Reciprocating
4	1.0	Cessna	NaN

6.3 Data Cleaning

Now that we have identified the columns we will be using, let us start on Data Preparation by looking for duplicates and missing values (and filling them)

```
In [13]: # Let us start by looking for duplicates
aviation_data.duplicated().sum()      #There is 79134 duplicates noted.

# Let us remove duplicates and save using inplace = True
aviation_data.drop_duplicates(inplace = True)

# Let us confirm that duplicates have been dropped
aviation_data.duplicated().sum()
```

Out[13]: 0

All Duplicates have been dropped

```
In [14]: # Let us now look for missing values and see how to fill them
aviation_data.isna().sum() # There are several missing values in each category
```

```
Out[14]: Total.Fatal.Injuries    1677
Make             12
Engine.Type     1649
dtype: int64
```

```
In [15]: # Our categories are both numerical (Total.Fatal.Injuries) and Non-numerical (Injury Se
# let us start with Total Fatal Injuries.
# we will first get the measures of central tendency to decide what to fill them with
aviation_data['Total.Fatal.Injuries'].agg(['mean', 'median'])
```

```
Out[15]: mean      2.326657
median    0.000000
Name: Total.Fatal.Injuries, dtype: float64
```

```
In [16]: # The mean of our Total Fatal Injuries is 2.3
# We will fill the missing values in this category with this mean
aviation_data['Total.Fatal.Injuries'].fillna(aviation_data['Total.Fatal.Injuries'].mean)
```

```
In [17]: # Let us confirm that our missing values in Total Fatal Injuries are done
aviation_data['Total.Fatal.Injuries'].isna().sum()
```

Out[17]: 0

```
In [18]: # Now let us consider dropping our non numerical columns instead of filling with mode a:
aviation_data.dropna(inplace = True)
```

```
In [19]: # Let us confirm that all our missing values in the non numerical columns are dropped.
aviation_data.isna().sum()
```

```
Out[19]: Total.Fatal.Injuries    0
Make             0
Engine.Type     0
dtype: int64
```

We have dropped all our missing data in our columns

In [20]: `#Let us check its effects on the total number of samples
len(aviation_data)`

Out[20]: 10509

Due to the dropna () our samples have dropped from 90348 to 11090

In [21]: `# Let us finally check for any data that is duplicated due to error in writing values
Let us start with the Makes columns

aviation_data['Make'].value_counts() # There is error in entry e.g: BOEING and boeing

aviation_data['Make'] = aviation_data['Make'].str.strip().str.title()

Get counts per manufacturer
make_counts = aviation_data['Make'].value_counts().reset_index()

Rename columns for clarity
make_counts.columns = ['Make', 'Count']`

In [22]: `# Let us confirm that we have normalized our outputs
aviation_data['Make'].value_counts()`

Out[22]:

Cessna	86
Boeing	74
Beech	73
Piper	50
McDonnell Douglas	45
	..
Gillette	1
Worley Dennis	1
Dennis L. Fetter	1
Chistov	1
Viking	1

Name: Make, Length: 7135, dtype: int64

In [23]: `# Let us also confirm that the value counts in our engine types are not repeated
aviation_data['Engine.Type'].value_counts() # Fortunately, these values were not repeated`

Out[23]:

Reciprocating	8633
Turbo Prop	537
Unknown	458
Turbo Shaft	362
Turbo Fan	283
Turbo Jet	202
None	16
Electric	10
LR	2
NONE	2
Geared Turbofan	2
UNK	1
Hybrid Rocket	1

Name: Engine.Type, dtype: int64

We have managed to clean our data. Now we start the analysis

6.4 Data Analysis

In [24]: `aviation_data.columns`

```
Out[24]: Index(['Total.Fatal.Injuries', 'Make', 'Engine.Type'], dtype='object')
```

```
In [25]: # So that we can determine which airplane makes to suggest to our company, we will compute the total fatal injuries by make.

fatal_by_make = aviation_data.groupby('Make', as_index=False)[['Total.Fatal.Injuries']].sum()

print(fatal_by_make.sort_values(by='Total.Fatal.Injuries', ascending=False))

      Make  Total.Fatal.Injuries
756    Boeing        3636.306629
4209  McDonnell Douglas       1037.633286
1831      Douglas        937.633286
156    Airbus Industrie       673.979971
151      Airbus        449.326657
...
4374      ...          ...
4375    Milentz        0.000000
4375    Miles Atwood        0.000000
4376      Mileski        0.000000
4377    Milholland        0.000000
7134  Zwicker Murray R        0.000000

[7135 rows x 2 columns]
```

```
In [28]: import seaborn as sns
import matplotlib.pyplot as plt

# --- Clean 'Make' column ---
aviation_data['Make'] = aviation_data['Make'].fillna('').str.strip().str.title()

# Remove empty Make values
aviation_data = aviation_data[aviation_data['Make'] != '']

# --- Group and sum total fatal injuries ---
fatal_by_make = aviation_data.groupby('Make', as_index=False)[['Total.Fatal.Injuries']].sum()

# --- Sort ascending (least fatalities first) and select top 10 ---
least_fatal_10 = fatal_by_make.sort_values(by='Total.Fatal.Injuries', ascending=True).head(10)


# --- Plot ---
plt.figure(figsize=(10, 6))
sns.barplot(
    data=least_fatal_10,
    y='Make',
    x='Total.Fatal.Injuries',
    palette='Greens'  # green = safer / fewer fatalities
)

plt.title('Top 10 Aircraft Makes with the Least Total Fatal Injuries', fontsize=14)
plt.xlabel('Total Fatal Injuries', fontsize=12)
plt.ylabel('Aircraft Make', fontsize=12)

# Add value labels
for index, value in enumerate(least_fatal_10['Total.Fatal.Injuries']):
    plt.text(value, index, f'{int(value)}', va='center', ha='left', fontsize=10)

plt.tight_layout()
plt.show()
```

Top 10 Aircraft Makes with the Least Total Fatal Injuries



```
In [29]: # Makes with the highest Total Fatal Injuries

# --- Clean 'Make' column ---
aviation_data['Make'] = aviation_data['Make'].fillna(' ').str.strip().str.title()


# --- Group and sum total fatal injuries ---
fatal_by_make = aviation_data.groupby('Make', as_index=False)[['Total.Fatal.Injuries']].sum()

# --- Sort descending (most fatalities first) and select top 10 ---
most_fatal_10 = fatal_by_make.sort_values(by='Total.Fatal.Injuries', ascending=False).head(10)

# --- Plot ---
plt.figure(figsize=(10, 6))
sns.barplot(
    data=most_fatal_10,
    y='Make',
    x='Total.Fatal.Injuries',
    palette='Reds_r' # darker = more fatalities
)
plt.title('Top 10 Aircraft Makes with the Most Total Fatal Injuries', fontsize=14)
plt.xlabel('Total Fatal Injuries', fontsize=12)
plt.ylabel('Aircraft Make', fontsize=12)

# Add value labels to bars
for index, value in enumerate(most_fatal_10['Total.Fatal.Injuries']):
    plt.text(value, index, f'{int(value)}', va='center', ha='left', fontsize=10)

plt.tight_layout()
plt.show()
```



From our results, the planes we surely would want to avoid due to safety are Boeing, Mc Douglas, Airbus, Beech, Cessna e.t.c (as shown above)

In [30]:

```

import seaborn as sns
import matplotlib.pyplot as plt

# --- Clean text columns ---
aviation_data['Make'] = aviation_data['Make'].fillna('').str.strip().str.title()
aviation_data['Engine.Type'] = aviation_data['Engine.Type'].fillna('').str.strip().str.title()

# Remove empty Make values
aviation_data = aviation_data[aviation_data['Make'] != '']

# --- Aggregate total fatal injuries by Make ---
fatal_by_make = (
    aviation_data.groupby('Make', as_index=False)[['Total.Fatal.Injuries']].sum()
)

# --- Select 30 makes with the least total fatal injuries ---
least_fatal_30 = fatal_by_make.nsmallest(30, 'Total.Fatal.Injuries')['Make']

# --- Filter original data for these 30 makes ---
filtered_data = aviation_data[aviation_data['Make'].isin(least_fatal_30)]

# Remove empty engine types
filtered_data = filtered_data[filtered_data['Engine.Type'] != '']

# --- Group by Make and Engine.Type for plotting ---
make_engine_counts = (
    filtered_data.groupby(['Make', 'Engine.Type'])
    .size()
    .reset_index(name='Count')
)

# Sort for cleaner plotting

```


```

make_engine_counts = make_engine_counts.sort_values(by='Count', ascending=False)

# --- Plot ---
plt.figure(figsize=(12, 10))
sns.barplot(
    data=make_engine_counts,
    x='Count',
    y='Make',
    hue='Engine.Type',
    palette='Set2'
)

plt.title('30 Aircraft Makes with Least Total Fatal Injuries by Engine Type', fontsize=14)
plt.xlabel('Number of Aircraft / Records', fontsize=12)
plt.ylabel('Aircraft Make', fontsize=12)
plt.legend(title='Engine Type', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

```



Reciprocating engines are good for long distance.

Turboprop are good for short distance.

Turbojet are good for medium distance

6.5 Results and Recommendations

Based on our results above, we would suggest that our company procures the 2007 Savage Air Lc for short distance flights and the other several suggested makes (as per chart) for long distance flights.